

Rewrite Rules for Relational Algebra with Applications to Query Optimization

Dirk Van Gucht¹

¹Indiana University

Relational Algebra Expressions

- The set of RA expressions can be recursively defined
- We will use the following notations:
 - E denotes an RA expression and \mathbf{A}_E denotes its schema (i.e., set of attributes)
 - F denotes an RA expression and \mathbf{A}_F denotes its schema (i.e., set of attributes)
 - C is a condition and \mathbf{A}_C denotes the set of attributes that occur in C
 - L denotes an attribute list and \mathbf{A}_L denotes the set of attributes in L

Relational Algebra (Recursive definition)

R $(A : \mathbf{a})$	with R a relation with A an attribute and \mathbf{a} a constant
$E \cup F$	with $\mathbf{A}_E = \mathbf{A}_F$
$E \cap F$	with $\mathbf{A}_E = \mathbf{A}_F$
$E - F$	with $\mathbf{A}_E = \mathbf{A}_F$
$\sigma_C(E)$	with $\mathbf{A}_C \subseteq \mathbf{A}_E$
$\pi_L(E)$	with $\mathbf{A}_L \subseteq \mathbf{A}_E$
$E \times F$	with $\mathbf{A}_E \cap \mathbf{A}_F = \emptyset$
$E \bowtie_C F$ $E \bowtie F$	with $\mathbf{A}_E \cap \mathbf{A}_F = \emptyset$ and $\mathbf{A}_C \subseteq (\mathbf{A}_E \cup \mathbf{A}_F)$
$E \ltimes F$ $E \overline{\ltimes} F$	

Conditions (Recursive definition)

The set of conditions can be recursively defined as follows:

$A \theta a$	with A an attribute, a a constant, and θ one of $=, \neq, <, \leq, >, \geq$
$A \theta B$	with A and B attributes and θ one of $=, \neq, <, \leq, >, \geq$
$C_1 \wedge C_2$	with C_1 and C_2 conditions
$C_1 \vee C_2$	with C_1 and C_2 conditions
$\neg C$	with C a condition
(C)	with C a condition

Query optimization overview

- Recall that SQL queries can be translated into equivalent RA expressions
- The benefit of this translation is that the declaratively specified SQL queries are transformed into procedurally specified queries (expressions)
- Nonetheless, these RA expressions can be inefficient to evaluate
- Rewriting these RA expressions can significantly improve this efficiency

Query optimization and SQL

- Recall that SQL can be used as a language to express RA expressions in close correspondence with RA's syntax
- Consequently, the principles for optimizing RA expressions can be applied to optimize SQL queries
- The translation algorithm from SQL to RA can be extended to incorporate optimization techniques developed for RA
- This is often a technique to improve the efficiency of SQL queries

How rewrite rules applied?

- Rewrite rules are expressed as **set equalities** between RA expressions. So they take the form

$$E = F$$

- Due to the property of and equality, a rewrite rule can be applied in two directions:
 - 1 From left to right: rewrite (replace) E by F
 - 2 From right to left: rewrite (replace) F by E
- For example,

$$A \cap (B \cup C) = (A \cap B) \cup (A \cap C) \quad (\text{Distribution of } \cap \text{ over } \cup)$$

- In applying this rule, it is sometimes convenient to replace $A \cap (B \cup C)$ by $(A \cap B) \cup (A \cap C)$, while, at other times, it is useful to do this in the other direction.

How is the correctness of a rewrite rule established?

- This is often done by using the proof techniques of
Predicate Logic
- For example, prove that $A \cap (B \cup C) = (A \cap B) \cup (A \cap C)$.
- We need to prove that
 - 1 $A \cap (B \cup C) \subseteq (A \cap B) \cup (A \cap C)$, and
 - 2 $(A \cap B) \cup (A \cap C) \subseteq A \cap (B \cup C)$
- To prove (1), take an element $x \in A \cap (B \cup C)$. This means that x is in A and x is in $B \cup C$.
Thus, x is in A and (a) x is in B or (b) x is in C .
If case (a) applies, then x is in $A \cap B$ and thus x is also $(A \cap B) \cup (A \cap C)$.
If case (b) applies, then x is in $A \cap C$ and thus x is also in $(A \cap B) \cup (A \cap C)$.
- The proof of (2) can be done in a similar way

Rewrite rules for the set operations \cup , \cap , and $-$

In this table, E , F , and G denote RA expressions with the same schemas

$E - (E - F) = E \cap F$	Double complementation of F relative to E When $E \supseteq F$
$E - (E - F) = F$	
$E - (F \cap G) = (E - F) \cup (E - G)$	Relativized De Morgan for \cap
$E - (F \cup G) = (E - F) \cap (E - G)$	Relativized De Morgan for \cup
$E \cap (F \cup G) = (E \cap F) \cup (E \cap G)$	Distribution of \cap over \cup
$E \cup (F \cap G) = (E \cup F) \cap (E \cup G)$	Distribution of \cup over \cap
$E \cap E = E$	Idempotence of \cap
$E \cup E = E$	Idempotence of \cup

Rewrite rules for the operations \cup , \cap , and $-$

In this table, E , F , and G denote RA expressions with the same schemas

$E \cap F$	$=$	$F \cap E$	Commutativity of \cap
$E \cup F$	$=$	$F \cup E$	Commutativity of \cup
$E \cap (F \cap G)$	$=$	$(E \cap F) \cap G$	Associativity of \cap
$E \cup (F \cup G)$	$=$	$(E \cup F) \cup G$	Associativity of \cup
$E \cap (E \cup F)$	$=$	E	Absorption for \cap
$E \cup (E \cap F)$	$=$	E	Absorption for \cup
$E \cap (F - E)$	$=$	\emptyset	Relativized contradiction for \cap
$E \cup (F - E)$	$=$	$E \cup F$	Relativized tautology for \cup
$E \cup \emptyset$	$=$	E	Identity for \cup
$E \cap \emptyset$	$=$	\emptyset	Domination for \cap

Rewrite rules for conditions

- Boolean conditions occur in both SQL and RA
 - In SQL, in the WHERE clause
 - In RA, in the selection operator σ
- Each of the rewrite rules for \cup , \cap , and $-$ has a corresponding **logical equivalence** between conditions involving **OR**, **AND**, and **NOT** in SQL, and involving \vee , \wedge , and \neg in RA.
- For example,

$$E_1 \cap (E_2 \cup E_3) = (E_1 \cap E_2) \cup (E_1 \cap E_3)$$

In SQL : $C_1 \text{ AND } (C_1 \text{ OR } C_3) \leftrightarrow (C_1 \text{ AND } C_2) \text{ OR } (C_1 \text{ AND } C_3)$

In RA : $C_1 \wedge (C_2 \vee C_3) \leftrightarrow (C_1 \wedge C_2) \vee (C_1 \wedge C_3)$

Rewrite rules for the selection operator σ

$$\sigma_{C_1 \wedge C_2}(E) = \sigma_{C_1}(\sigma_{C_2}(E)) \quad \text{Cascading selections}$$

$$= \sigma_{C_2}(\sigma_{C_1}(E)) \quad \text{Commutativity of selections}$$

$$\sigma_{C_1 \wedge C_2}(E) = \sigma_{C_1}(E) \cap \sigma_{C_2}(E) \quad \text{Boolean decomposition of } \wedge$$

$$\sigma_{C_1 \vee C_2}(E) = \sigma_{C_1}(E) \cup \sigma_{C_2}(E) \quad \text{Boolean decomposition of } \vee$$

$$\sigma_{\neg C}(E) = E - \sigma_C(E) \quad \text{Boolean decomposition of } \neg$$

Rewrite rules for the selection operator σ (Distribution rules)

$$\sigma_C(E \cup F) = \sigma_C(E) \cup \sigma_C(F)$$

$$\begin{aligned}\sigma_C(E \cap F) &= \sigma_C(E) \cap \sigma_C(F) \\ &= \sigma_C(E) \cap F \\ &= E \cap \sigma_C(F)\end{aligned}$$

$$\begin{aligned}\sigma_C(E - F) &= \sigma_C(E) - \sigma_C(F) \\ &= \sigma_C(E) - F\end{aligned}$$

Rewrite rules for interactions between selection σ and join \bowtie operations

The rules of **pushing down selections over joins** are frequently used since they can substantially improve the efficiency of evaluating expressions!

$\sigma_C(E \times F)$	$=$	$E \bowtie_C F$	Definition of \bowtie_C
$\sigma_C(E \bowtie F)$	$=$	$\sigma_C(E) \bowtie F$	when $\mathbf{A}_C \subseteq \mathbf{A}_E$
	$=$	$E \bowtie \sigma_C(F)$	when $\mathbf{A}_C \subseteq \mathbf{A}_F$
	$=$	$\sigma_C(E) \bowtie \sigma_C(F)$	when $\mathbf{A}_C \subseteq \mathbf{A}_E \cap \mathbf{A}_F$
$\sigma_{C_1}(E \bowtie_{C_2} F)$	$=$	$\sigma_{C_1}(E) \bowtie_{C_2} F$	when $\mathbf{A}_{C_1} \subseteq \mathbf{A}_E$
	$=$	$E \bowtie_{C_2} \sigma_{C_1}(F)$	when $\mathbf{A}_{C_1} \subseteq \mathbf{A}_F$
	$=$	$E \bowtie_{C_1 \wedge C_2} F$	
$\sigma_C(E \bowtie F)$	$=$	$\sigma_C(E) \bowtie F$	
$\sigma_C(E \overline{\bowtie} F)$	$=$	$\sigma_C(E) \overline{\bowtie} F$	

Rewrite rules for the interactions of projections π and the set operations \cup , \cap , and $-$

Be careful with pushing down (i.e., distributing) projections π over the set operations \cup , \cap and $-$.

Recall the lecture on the translation algorithm from SQL to RA.

$$\pi_L(E \cup F) = \pi_L(E) \cup \pi_L(F)$$

$$\pi_L(E \cap F) \subseteq \pi_L(E) \cap \pi_L(F)$$

$$\pi_L(E - F) \supseteq \pi_L(E) - \pi_L(F)$$

Rewrite rules for the projection operator

In general, it is **not** the case that $\pi_L(E \cap F) = \pi_L(E) \cap \pi_L(F)$.

The following is a counter example:

R		S	
A	B	A	B
a	b	a	c

In this case,

$$\pi_A(R \cap S) = \pi_A(\emptyset) = \emptyset$$

but

$$\pi_A(R) \cap \pi_A(S) = \{a\} \cap \{a\} = \{a\}$$

Rewrite rules for the projection operator π

In general, it is **not** the case that $\pi_L(E - F) = \pi_L(E) - \pi_L(F)$.

The following is a counter example:

R		S	
A	B	A	B
a	b	a	c

$$\pi_A(R - S) = \pi_A(\{(a, b)\}) = \{a\}$$

but

$$\pi_A(R) - \pi_A(S) = \{a\} - \{a\} = \emptyset$$

Rewrite rules for the interactions of π and σ

Projection and commute

$$\pi_L(\sigma_C(E)) = \sigma_C(\pi_L(E)) \quad \text{when } \mathbf{A}_L \subseteq \mathbf{A}_C$$

It is often not clear in which direction to apply this rule since both $\sigma_C(E)$ and $\pi_L(E)$ are space reducing

And it is not always clear which of these space reductions is best. This depends on the data.

Elimination rules for π

$$\pi_L(E) = E$$

when the schema of E corresponds precisely with L

$$\pi_{B,A}(E) \neq E$$

if the schema of E is (A, B)
projection acts as a permutation

$$\pi_{L_1}(\pi_{L_2}(E)) = \pi_{L_1}(E)$$

Rewrite rules for the interactions of projections π and joins \bowtie

- Observe that we can expect difficulties when we look for rules involving π and joins (\bowtie_C , \bowtie , \ltimes , and $\overline{\bowtie}$)
- Recall that $E \cap F = E \bowtie F$
- We know that, in general, $\pi_L(E \cap F)$ is not the same as $\pi_L(E) \cap \pi_L(F)$
- Therefore, we conclude that, in general, π does not distributes over joins

Rewrite rules for the interactions of projections π and joins \bowtie

We have the following important rule:

$$\pi_L(E \bowtie_C F) = \pi_L(\pi_{\mathbf{A}_{L_E}}(E) \bowtie_C \pi_{\mathbf{A}_{L_F}}(F))$$

with $\mathbf{A}_{L_E} = \mathbf{A}_E \cap (\mathbf{A}_L \cup \mathbf{A}_C)$ and

with $\mathbf{A}_{L_F} = \mathbf{A}_F \cap (\mathbf{A}_L \cup \mathbf{A}_C)$

This rule permits us to project-out (eliminate) each attribute from E (or from F) that does not appear in both

- 1 the projection list L and
- 2 the join condition C

This rule is called the **attribute elimination** rule or the rule of **pushing projections down over joins**.

Rewrite rules for π and joins \bowtie (Example)

Consider

$$\pi_{a,i}(E \bowtie_{a=g} F)$$

and assume that $\mathbf{A}_E = \{a, b\}$ and $\mathbf{A}_F = \{g, h, i\}$

By the attribute elimination rule, we have that

$$\begin{aligned}\pi_{a,i}(E \bowtie_{a=g} F) &= \pi_{a,i}(\pi_{\mathbf{A}_{L_E}}(E) \bowtie_{a=g} \pi_{\mathbf{A}_{L_F}}(F)) \\ &= \pi_{a,i}(\pi_a(E) \bowtie_{a=g} \pi_{g,i}(F))\end{aligned}$$

since

$$\begin{aligned}\mathbf{A}_{L_E} &= \mathbf{A}_E \cap (\mathbf{A}_L \cup \mathbf{A}_C) = \{a, b\} \cap (\{a, i\} \cup \{a, g\}) = \{a\} \\ \mathbf{A}_{L_F} &= \mathbf{A}_F \cap (\mathbf{A}_L \cup \mathbf{A}_C) = \{g, h, i\} \cap (\{a, i\} \cup \{a, g\}) = \{g, i\}\end{aligned}$$

So attributes b and h were eliminated since they not appear in both the projection list and the join condition

Rewrite rule for joins \bowtie_C and natural join \bowtie

Assume that E and F have overlapping attributes B_1, \dots, B_k .

Then

$$\pi_L(E \bowtie_{E.B_1=F.B_1 \wedge \dots \wedge E.B_k=F.B_k} F) = \pi_L(E \bowtie F)$$

Notice that we permit both $E.B_i$ and $F.B_i$ to simultaneously occur in L .

Of course, then the columns $E.B_i$ and $F.B_i$ in $\pi_L(R \bowtie S)$ are identical and therefore, in essence, one of them is redundant

Regular Semi-joins and Anti-semijoins

- Since regular semi-joins are combination of joins and projection, all the laws relative these operations apply.

$$E \ltimes F = \pi_{\mathbf{A}_E}(E \bowtie F)$$

- We also have the following important rewrite rule

$$E \ltimes F = E \bowtie \pi_{\mathbf{A}_E \cap \mathbf{A}_F}(F)$$

- For anti-joins, we have

$$E \bar{\ltimes} F = E - (E \ltimes F)$$

Thus all the laws relative to the interactions of set-difference and semi-joins apply

Example 1

- “Find the name of each student who is enrolled in course 2003.”

```
SELECT  s.sname
FROM    Student s, Enroll e
WHERE   s.sid = e.sid AND e.cno = 2003
```

- This query is equivalent with the RA expression

$$\pi_{sname}(\sigma_{S.sid=E.sid \wedge cno=2003}(S \times E))$$

where S denotes Student and E denotes Enroll

Example 1 (Optimization)

- “Find the name of each student who is enrolled in course 2003.”
- In RA,

$$\pi_{sname}(\sigma_{S.sid=E.sid \wedge cno=2003}(S \times E))$$

- Optimization:

$$\begin{aligned} & \pi_{sname}(\sigma_{S.sid=E.sid \wedge cno=2003}(S \times E)) & = \\ & \pi_{sname}(\sigma_{S.sid=E.sid}(\sigma_{cno=2003}(S \times E))) & = \\ & \pi_{sname}(\sigma_{S.sid=E.sid}(S \times \sigma_{cno=2003}(E))) & = \\ & \pi_{sname}(S \bowtie_{S.sid=E.sid} \sigma_{cno=2003}(E)) & = \\ & \pi_{sname}(\pi_{sname, S.sid}(S) \bowtie_{S.sid=E.sid} \pi_{E.sid}(\sigma_{cno=2003}(E))) & = \\ & \pi_{sname}(\pi_{sname, sid}(S) \bowtie \pi_{sid}(\sigma_{cno=2003}(E))) & = \\ & \pi_{sname}(\pi_{sname, sid}(S) \bowtie \pi_{sid}(\sigma_{cno=2003}(E))) \end{aligned}$$

Example 1 (Revisited)

- “Find the name of each student who is enrolled in course 2003.”

```
SELECT  DISTINCT s.sname
FROM    Student s, Enroll e
WHERE   s.sid = e.sid AND e.cno = 2003
```

- The SQL-to-RA translation algorithm could have produced the RA expression

$$\pi_{sname}(S \bowtie \sigma_{cno=2003}(E))$$

- Optimization:

$$\begin{aligned}\pi_{sname}(S \bowtie \sigma_{cno=2003}(E)) &= \\ \pi_{sname}(\pi_{sname,sid}(S) \bowtie \pi_{sid}(\sigma_{cno=2003}(E))) &= \\ \pi_{sname}(\pi_{sname,sid}(S) \ltimes \pi_{sid}(\sigma_{cno=2003}(E)))\end{aligned}$$

Example 1 (Revisited)

- “Find the name of each student who is enrolled in course 2003.”

```
SELECT  DISTINCT s.sname
FROM    Student s, Enroll e
WHERE   s.sid = e.sid AND e.cno = 2003
```

- The SQL-to-RA translation algorithm could have produced the RA expression

$$\pi_{sname}(\sigma_{cno=2003}(S \bowtie E))$$

- Optimization:

$$\begin{aligned}\pi_{sname}(\sigma_{cno=2003}(S \bowtie E)) &= \\ \pi_{sname}(S \bowtie \sigma_{cno=2003}(E)) &= \\ \pi_{sname}(\pi_{sname,sid}(S) \bowtie \pi_{sid}(\sigma_{cno=2003}(E))) &= \\ \pi_{sname}(\pi_{sname,sid}(S) \ltimes \pi_{sid}(\sigma_{cno=2003}(E)))\end{aligned}$$

Example 2

- “Find the sid of each student who is enrolled in at least one CS course.”

```
SELECT  s.sid
FROM    Student s
WHERE   EXISTS (SELECT 1
                  FROM    Enroll e, Course c
                  WHERE   s.sid = e.sid AND e.cno = c.cno AND
                          dept = 'CS' )
```

- This query is equivalent with the RA expression

$$\pi_{S.sid}(\sigma_{S.sid=E.sid \wedge E.cno=C.cno \wedge dept='CS'}(S \times E \times C))$$

where S denotes Student, E denotes Enroll, and C denotes Course

Example 2

- “Find the sid of each student who is enrolled in at least one CS course.”
- Optimization:

$$\begin{aligned}
 & \pi_{S.sid}(\sigma_{S.sid=E.sid \wedge E.cno=C.cno \wedge dept='CS'}(S \times E \times C)) & = \\
 & \pi_{S.sid}(\sigma_{S.sid=E.sid}(\sigma_{E.cno=C.cno}(\sigma_{dept='CS'}(S \times E \times C)))) & = \\
 & \pi_{S.sid}(S \bowtie_{S.sid=E.sid} (E \bowtie_{E.cno=C.cno} \sigma_{dept='CS'}(C))) & = \\
 & \pi_{sid}(S \bowtie (E \bowtie \sigma_{dept='CS'}(C))) & = \\
 & \pi_{sid}(\pi_{sid}(S) \bowtie \pi_{sid}(\pi_{sid,cno}(E) \bowtie \pi_{cno}(\sigma_{dept='CS'}(C)))) & = \\
 & \pi_{sid}(S) \bowtie \pi_{sid}(\pi_{sid,cno}(E) \bowtie (\pi_{cno}(\sigma_{dept='CS'}(C)))) & = \\
 & \pi_{sid}(S) \cap \pi_{sid}(\pi_{sid,cno}(E) \bowtie (\pi_{cno}(\sigma_{dept='CS'}(C)))) & = \\
 & \pi_{sid}(\pi_{sid,cno}(E) \bowtie \pi_{cno}(\sigma_{dept='CS'}(C))) & =
 \end{aligned}$$

- The last equality follows since

$$\pi_{sid}(\pi_{sid,cno}(E) \bowtie (\pi_{cno}(\sigma_{dept='CS'}(C)))) \subseteq \pi_{sid}(S)$$

This is because *sid* is a **foreign key** in Enroll referencing the primary key *sid* in Student.

Example 2 (Revisited)

- “Find the sid of each student who is enrolled in at least one CS course.”

```
SELECT  s.sid
FROM    Student s
WHERE   EXISTS (SELECT 1
                  FROM    Enroll e, Course c
                  WHERE   s.sid = e.sid AND e.cno = c.cno AND
                          dept = 'CS' )
```

- The SQL-to-RA algorithm could have produced the RA expression:

$$\pi_{sid}(S \bowtie (E \bowtie \sigma_{dept='CS'}(C)))$$

- Optimization:

$$\begin{aligned} \pi_{sid}(S \bowtie (E \bowtie \sigma_{dept='CS'}(C))) &= \\ \pi_{sid}(\pi_{sid}(S) \bowtie \pi_{sid}(\pi_{sid,cno}(E) \bowtie \pi_{cno}(\sigma_{dept='CS'}(C)))) &= \\ \pi_{sid}(S) \bowtie \pi_{sid}(\pi_{sid,cno}(E) \bowtie (\pi_{cno}(\sigma_{dept='CS'}(C)))) &= \\ \pi_{sid}(S) \cap \pi_{sid}(\pi_{sid,cno}(E) \bowtie (\pi_{cno}(\sigma_{dept='CS'}(C)))) &= \\ \pi_{sid}(\pi_{sid,cno}(E) \bowtie \pi_{cno}(\sigma_{dept='CS'}(C))) & \end{aligned}$$

Example 3

- “Find the sid of each student who takes a course also taken by student with sid = s100.”

```
SELECT  DISTINCT e1.sid
FROM    Enroll e1, Enroll e2
WHERE   e1.cno = e2.cno AND e2.sid = 's100'
```

- The SQL-to-RA algorithm could have produced the RA expression:

$$\pi_{E_1.sid}(E_1 \bowtie_{E_1.cno=E_2.cno} (\sigma_{sid='s100'}(E_2)))$$

- Optimization:

$$\begin{array}{ll} (1) & \pi_{E_1.sid}(E_1 \bowtie_{E_1.cno=E_2.cno} (\sigma_{sid='s100'}(E_2))) & = \\ & \pi_{E_1.sid}(\pi_{E_1.sid, E_1.cno}(E_1) \bowtie_{E_1.cno=E_2.cno} \pi_{E_2.cno}(\sigma_{E_2.sid='s100'}(E_2))) & = \\ & \pi_{sid}(\pi_{sid, cno}(E_1) \bowtie \pi_{cno}(\sigma_{sid='s100'}(E_2))) & = \\ (2) & \pi_{sid}(\pi_{sid, cno}(E_1) \bowtie \pi_{cno}(\sigma_{sid='s100'}(E_2))) & = \end{array}$$

- Expression (1) is $O(|Enroll|^2)$ but expression (2) is just $O(|Enroll|)$
- Optimization results in order of magnitude improvement

Constraints and Optimization

- The presence of constraints introduces other optimization opportunities
- We consider primary key and foreign key constraints
- The optimizations apply to interactions of projections and the intersection, join and set difference operations.

Primary keys and distribution of projection over intersection and set difference

- Consider a **primary key** \mathbf{K} in a relation R
- Then \mathbf{K} uniquely determines the values of the tuples at each subset \mathbf{A} of the set of attributes in R
- So we have $R(\mathbf{K}, \mathbf{A}, \dots)$ with \mathbf{K} the primary key
- We write

$$\mathbf{K} \rightarrow \mathbf{A}$$

- We have the following rewrite rules

$$\pi_{\mathbf{K}}(\pi_{\mathbf{K}, \mathbf{A}}(E_1) \cap \pi_{\mathbf{K}, \mathbf{A}}(E_2)) = \pi_{\mathbf{K}}(E_1) \cap \pi_{\mathbf{K}}(E_2)$$

$$\pi_{\mathbf{K}}(\pi_{\mathbf{K}, \mathbf{A}}(E_1) - \pi_{\mathbf{K}, \mathbf{A}}(E_2)) = \pi_{\mathbf{K}}(E_1) - \pi_{\mathbf{K}}(E_2)$$

Primary keys and distribution of projection over joins, intersection, and set difference (Generalization)

- Consider the functional constraint

$$\mathbf{K} \rightarrow \mathbf{A}$$

- Then we have the following rewrite rules that **eliminate the attributes in \mathbf{A}**

$$\pi_{\mathbf{L}}(\pi_{\mathbf{K}, \mathbf{A}, \mathbf{B}_1}(E_1) \bowtie \pi_{\mathbf{K}, \mathbf{A}, \mathbf{B}_2}(E_2)) = \pi_{\mathbf{L}}(\pi_{\mathbf{K}, \mathbf{B}_1}(E_1) \bowtie \pi_{\mathbf{K}, \mathbf{B}_2}(E_2))$$

$$\pi_{\mathbf{L}}(\pi_{\mathbf{K}, \mathbf{A}, \mathbf{B}}(E_1) \cap \pi_{\mathbf{K}, \mathbf{A}, \mathbf{B}}(E_2)) = \pi_{\mathbf{L}}(\pi_{\mathbf{K}, \mathbf{B}}(E_1) \cap \pi_{\mathbf{K}, \mathbf{B}}(E_2))$$

$$\pi_{\mathbf{L}}(\pi_{\mathbf{K}, \mathbf{A}, \mathbf{B}}(E_1) - \pi_{\mathbf{K}, \mathbf{A}, \mathbf{B}}(E_2)) = \pi_{\mathbf{L}}(\pi_{\mathbf{K}, \mathbf{B}}(E_1) - \pi_{\mathbf{K}, \mathbf{B}}(E_2))$$

Foreign keys and natural joins

- $R(\mathbf{F}, \mathbf{B})$ and $S(\mathbf{K}, \mathbf{A})$
- Let \mathbf{F} be a **foreign key** in relation R referencing the primary key \mathbf{K} of relation S
- Let $L \subseteq \mathbf{F} \cup \mathbf{B}$ and $\mathbf{B} \cap \mathbf{A} = \emptyset$
- We then have the following rewrite rule

$$\pi_L(R \bowtie S) = \pi_L(S \bowtie R) = \pi_L(R)$$

- So the relation S has been eliminated from this expression
- Indeed, we have

$$\begin{aligned}\pi_L(R \bowtie S) &= \pi_L(\pi_{L \cup \mathbf{F}}(R) \bowtie \pi_{\mathbf{F}}(S)) && \text{pushing } \pi \text{ over } \bowtie \\ &= \pi_L(\pi_{L \cup \mathbf{F}}(R)) && \mathbf{F} \text{ is a foreign key thus} \\ &&& \text{each tuple in } \pi_{L \cup \mathbf{F}}(R) \\ &&& \text{survives the join} \\ &= \pi_L(R)\end{aligned}$$

Foreign keys and natural joins (Example)

- *Enroll*(**sid**, *cno*, *grade*) and *Student*(**sid**, *sname*, *byear*)
- *sid* in *Enroll* is a FK referencing the primary key *sid* in *Student*

$$\pi_{sid, cno}(Enroll \bowtie Student) = \pi_{sid, cno}(Student \bowtie Enroll) = \pi_{sid, cno}(E)$$

$$\pi_{sid}(Enroll \bowtie Student) = \pi_{sid}(Student \bowtie Enroll) = \pi_{sid}(E)$$

$$\pi_{cno}(Enroll \bowtie Student) = \pi_{cno}(Student \bowtie Enroll) = \pi_{cno}(E)$$

- The relation *Student* has been eliminated in these expressions. I.e., the join \bowtie does not need to be done

Example 4

- “Find the cno of each course in which no students are enrolled.”

```
SELECT  cno
FROM    Course
WHERE   cno NOT IN (SELECT cno
                    FROM    Enroll )
```

- After applying the translation algorithm, we get the SQL query

```
SELECT  cno
FROM    (SELECT cno, cname, dept
        FROM    Course c
        EXCEPT
        SELECT cno, cname, dept
        FROM    Course NATURAL JOIN Enroll) q
```

Example 4 (Constraints)

- “Find the *cno* of each course in which no students are enrolled.”

```
SELECT  cno
FROM    (SELECT cno, cname, dept
        FROM   Course
        EXCEPT
        SELECT cno, cname, dept
        FROM   Course NATURAL JOIN Enroll) q
```

- This gets translated to the RA expression

$$\begin{array}{lll} (1) & \pi_{cno}(\pi_{cno, cname, dept}(C) - \pi_{cno, cname, dept}(C \bowtie E)) & = \\ (2) & \pi_{cno}(C) - \pi_{cno}(C \bowtie E) & = \\ (3) & \pi_{cno}(C) - \pi_{cno}(E) & \end{array}$$

- The rule that takes (1) to (2) follows since *cno* is a **primary key** for Course
- The rule that takes (2) to (3) follows since *cno* is **foreign key** in Enroll referencing the primary key *cno* in Course

Example 5

“Find the sid of each student who is only enrolled in CS courses.”

Using the translation algorithm this becomes the SQL query

```
WITH          CS AS (SELECT * FROM Course WHERE dept = 'CS')
SELECT        q1.ssid
FROM          (SELECT s.sid AS ssid, s.sname
               FROM    Student s
               EXCEPT
               SELECT  q2.ssid, s.sname
               FROM    (SELECT s.sid AS ssid, s.sname, e.sid, e.cno
                       FROM    Student s NATURAL JOIN Enroll e
                       EXCEPT
                       (SELECT s.sid, s.sname, e.sid, e.cno
                        FROM    Student s CROSS JOIN
                               Enroll e NATURAL JOIN CS ) q2) q1
```

is correspond to the RA expression

$$\pi_{sid}(\pi_{sid,sname}(S) - \pi_{sid,sname}(\pi_{S.sid,sname,E.sid,cno}(S \bowtie E) - \pi_{S.sid,sname,E.sid,cno}(S \bowtie E \bowtie \sigma_{dept='CS'}(C))))$$

where S , E , and C denote Student, Enroll, and Course, respectively.

Example 5

“Find the sid of each student who is only enrolled in CS courses.”

$$\pi_{sid}(\pi_{sid,sname}(S) - \pi_{sid,sname}(\pi_{S.sid,sname,E.sid,cno}(S \bowtie E) - \pi_{S.sid,sname,E.sid,cno}(S \bowtie E \bowtie \sigma_{dept='CS'}(C))))$$

Because *sid* is a primary key of Student and we do not need *sname*, we can rewrite this expression to

$$\pi_{sid}(S) - \pi_{sid}(\pi_{S.sid,E.sid,cno}(S \bowtie E) - \pi_{S.sid,E.sid,cno}(S \bowtie E \bowtie \sigma_{dept='CS'}(C))))$$

Because *sid* is a FK in Enroll referencing the primary key *sid* in Student, we can rewrite this expression to

$$\pi_{sid}(S) - \pi_{sid}(\pi_{sid,cno}(E) - \pi_{sid,cno}(E \bowtie \sigma_{dept='CS'}(C))))$$

If Enroll has schema (*sid*, *cno*) then this is

$$\pi_{sid}(S) - \pi_{sid}(E - E \bowtie \pi_{cno}(\sigma_{dept='CS'}(C))))$$

Example 5

"Find the sid of each student who is only enrolled CS courses."

```
SELECT  s.sid
FROM    Student s
WHERE   NOT EXISTS (SELECT 1
                    FROM    Enroll e
                    WHERE e.sid = s.sid AND
                          e.cno NOT IN (SELECT c.cno
                                        FROM    Course c
                                        WHERE   c.dept = 'CS'))
```

is translated and optimized to

$$\pi_{sid}(S) - \pi_{sid}(E \leftarrow E \bowtie \pi_{cno}(\sigma_{dept='CS'}(C)))$$

Or, more succinctly, using the anti-semijoin

$$\pi_{sid}(S) - \pi_{sid}(E \overline{\bowtie} \pi_{cno}(\sigma_{dept='CS'}(C)))$$